

ClosetAI — *the style core.*

Turning Appendix A3 of my master document into a working generative module that anyone can use in one minute, with no account and no photos.

Tamara Muñoz Delgadillo · Universidad Iberoamericana
Submitted 9 September 2026 · Product live at **closetai.lat**

LIVE PAGE closetai.lat/core

GITHUB github.com/vestiamx-code/closetai

DEMO VIDEO closetai.lat/demo/semana-1

THE FEATURE IN ONE LINE. ClosetAI asks you to photograph ten garments before it gives you anything. That was the strongest objection my value proposition received, and it is what killed the other apps in this category. `/core` reverses the order: you write three sentences about how you like to dress, and you get a structured style core back — essence, principles, palette, silhouettes, what to avoid, and one rule of your own. No sign-up. No upload. If your text is vague, it says so instead of inventing a palette for you.

1 · Weekly Submission Packet — every required item

ITEM	REQUIRED	DELIVERED	STATUS
Live URL	Deployed page	<code>closetai.lat/core</code> — HTTP 200 with no session cookie, verified at submission time (§4)	✓
Build Discipline Packet	Complete before coding	All 11 sections (§2). Committed in <code>a4e9890</code> , 17 minutes before the first line of feature code in <code>e02b657</code> — git history proves the order	✓
UX image / mockup	Image-generated concept	Image-generated mockup + three-state wireframe (§3)	✓
Product Spec	Requirements + acceptance criteria	8 requirements, each with a checkable criterion (§2)	✓
Architecture Sketch	Data flow + components	Full flow from browser to Postgres, with the contract gate marked (§2)	✓
GitHub commits	Minimum 5	15 commits this week, all pushed and public (§5)	✓
Vercel deployments	Minimum 2	9 production deployments , all Ready (§6)	✓
Supabase evidence	Table / data evidence	Migration <code>005_core_outputs</code> + 10 real rows queried live from production (§7)	✓
Prompt log	Minimum 5 prompts	6 prompts , each with result and the human judgment applied (§8)	✓
Test evidence	Minimum 3 tests	3 live runs against production, plus 13 new automated tests, plus a full pre-submission verification against the live site (§9)	✓
Iteration Log	What changed after testing	Three entries: two badly written tests, the provider quota, and three red tests found the day of submission (§10)	✓
Demo video	2–3 minutes	2:14 , hosted on the product's own domain (§11)	✓
Human Decision Note	150–250 words	239 words (§12)	✓

Required features and coding tasks from the week overview

REQUIRED	WHERE IT IS
Intake form	<code>/core</code> , 40–1000 characters, live counter, validated before any model call is spent
Core extraction output card	Structured card: essence, principles, palette, silhouettes, what to avoid, one rule
Save button	Server Action <code>guardarNucleo</code> , which re-validates the payload server-side
Supabase table <code>core_outputs</code>	Migration <code>20260902000005</code> , applied to production, 7 rows
Dashboard preview	"Núcleos que ya se generaron" — the saved cores, below the form
Prompt library entry	<code>src/lib/ai/prompts/core.ts</code> , versioned like the project's other prompts
3 test runs	§9 — three different inputs against the live site, one vague on purpose

2 · Build Discipline Packet

Written before touching code. The commit that added this file (`a4e9890` , 20:16) precedes the commit that added the feature (`e02b657` , 20:33). The history is public and checkable.

Problem

For ClosetAI to be useful today, you have to photograph about ten garments. The value arrives **after** an hour of work. That is the strongest objection my value proposition received, and it is the same thing that killed the other apps in this category: they make you catalogue your closet, and only then give you something. Nobody can tell whether ClosetAI's taste is worth that hour before spending it.

User

Someone landing on closetai.lat for the first time, with no account and no photos. It also serves people who already have an account: their style core feeds the stylist from day one, instead of waiting for five outfit reactions to accumulate.

Success — what has to work by the end of the week

- `/core` loads in production, **with no session**, in a private window.
- I write in my own words how I like to dress and get back a **structured style core**, not a loose paragraph.
- I can save it and see it again on the same page.
- Every save lands in the `core_outputs` table in Supabase.
- Three runs with different inputs produce three different, believable cores.

UX concept

One screen, three states: empty form → generating → result card with a save button. Below it, what has already been saved. Mockup and wireframe in §3.

Scope cut — what looked nice and is deliberately not built

- **Accounts on `/core`**. Saving without a session is the whole point: asking someone to register before proving value is exactly the problem this feature attacks.
- **Editing a saved core**. You can generate again; field-by-field editing is another week's work.
- **Wiring the core into the stylist**. It is the obvious next step and that is precisely why it waits: first I have to prove the core that comes out is any good.
- **Photo upload on `/core`**. The entire point is that none is needed.

Product spec and acceptance criteria

#	REQUIREMENT	ACCEPTANCE CRITERION (CHECKABLE)
1	Public page	<code>GET /core</code> returns 200 with no session cookie
2	Intake form	Accepts 40–1000 characters; outside that range, a clear error in Spanish

#	REQUIREMENT	ACCEPTANCE CRITERION (CHECKABLE)
3	Core extraction	The model's response passes a zod contract before it is displayed; if it fails, an error is shown and nothing is saved
4	Structured card	Shows essence, principles, palette, silhouettes, what to avoid, and one rule
5	Save	Writes a row to <code>core_outputs</code> and the row appears in the list without a reload
6	Saved view	Shows the most recent saved cores
7	Cost logged	Every generation writes its cost to <code>api_costs</code>
8	Usage limit	One IP cannot generate more than 10 times per hour

Architecture

```

Browser (/core - no session)
  | form → Server Action generarNucleo()
  ▼
Server Action
  | validates the input with zod
  | checks the per-IP rate limit
  | calls Gemini with CORE_PROMPT (versioned)
  | validates the output with styleCoreSchema ← if this fails: error, nothing saved
  | logs the cost to api_costs
  |
  ▼ (the person decides to save)
Server Action guardarNucleo() ← re-validates; a Server Action is reachable by direct POST
  | inserts into core_outputs (admin client; RLS allows public read only)
  |
  ▼
Supabase Postgres · core_outputs
  | the public list SELECTs every column except `entrada`

```

Why a Server Action and not an API route: the Gemini key never leaves the server, and there is no public endpoint for anyone to hit directly.

Tech stack

TOOL	FOR WHAT	WHY THIS ONE
Next.js 16 · Server Actions	Page and server logic	Already the project's stack; the API key stays server-side
Gemini <code>flash</code>	Extracting the core	This is reasoning over text, not vision — the cheap model is enough
zod	Output contract	A model returns free text sooner or later; this is the fence
Supabase Postgres	<code>core_outputs</code> table	The project and its migrations already exist
Tailwind v4	Styling	Same design tokens as the rest of the app

DevOps

- Repo `vestiamx-code/closetai`, public. Branch `main`, automatic deploy.

- New migration `005_core_outputs.sql`, applied to production *before* deploying.
- No new environment variables: reuses `GEMINI_API_KEY` and the Supabase ones.
- GitHub Actions already runs lint, types, unit tests and build on every push.

Test plan

Automated — unit: the contract rejects a core with no `esencia`; the contract rejects more than 5 principles. End-to-end: `/core` loads with no session and shows the form; valid text produces a card with all its fields; the save button makes the core appear in the list.

Three manual runs against the live site, with different inputs, recording input / expected result / actual result / what failed / what I changed.

Coding agent prompt

Build the public page `/core` in the ClosetAI project (Next.js 16, App Router, Server Actions, Tailwind v4, Supabase).

Use the methodology in Appendix A3 of the master document as the source and turn it into a generative module: the user writes in her own words how she likes to dress, and the module returns a structured style core – essence, principles, palette, silhouettes, what to avoid, and one personal rule.


Requirements: the page is public, no session. The model's output goes through a zod contract before being displayed and is not saved if it does not comply. A save button that writes to the `core_outputs` table. Below it, the list of saved cores. Log the cost of each call to `api_costs` and limit to 10 generations per IP per hour. Write the prompt in `src/lib/ai/prompts/core.ts`, versioned, like the ones that already exist. Do not invent columns: write the migration first.

3 • UX image and wireframe

The mockup was image-generated from the concept above. The wireframe shows the three states the page actually has.



Image-generated mockup of the core card.



Three states: empty form → generating → result card with save.

4 • Live page

Verified at submission time, in a private window, against the custom domain — not localhost:

```
$ curl -s -o /dev/null -w "HTTP %{http_code} · %{time_total}s\n" https://closetai.lat/core
HTTP 200 · 3.39s
```

ClosetAI [Crear cuenta](#)

NUCLEO DE ESTILO

Antes de fotografiar una sola prenda.

Cuéntame cómo te gusta vestirte y destilo tu núcleo de estilo: tus principios, tu paleta, tus siluetas y una regla tuya. Sin cuenta, sin subir fotos, en menos de un minuto.

¿Cómo te gusta vestirte?
Escribelo como se lo contarías a una amiga. No hace falta que suene bonito.

Me gusta andar cómoda pero que se vea que me arreglé. Uso mucho negro y beige, jeans anchos y tenis. Odio la ropa muy entallada y los estampados grandes. Para la oficina le pongo un saco encima y ya.

Usar un ejemplo 0 / 1000

[Extraer mi núcleo](#)

Núcleos que ya se generaron
Solo se muestra el núcleo. Lo que cada persona escribió sobre sí misma no se publica nunca.

The empty state, in a private window with no session.

ClosetAI [Crear cuenta](#)

NUCLEO DE ESTILO

Antes de fotografiar una sola prenda.

Cuéntame cómo te gusta vestirte y destilo tu núcleo de estilo: tus principios, tu paleta, tus siluetas y una regla tuya. Sin cuenta, sin subir fotos, en menos de un minuto.

¿Cómo te gusta vestirte?
Escribelo como se lo contarías a una amiga. No hace falta que suene bonito.

Me gusta andar cómoda pero que se note que me arreglé. Uso mucho negro y beige, jeans anchos y tenis blancos. No me gusta la ropa entallada ni los estampados grandes. Para la oficina le pongo un saco encima y ya la hice.

Usar un ejemplo 228 / 1000

[Extraer mi núcleo](#)

TU NUCLEO DE ESTILO

Estilo casual y cómodo basado en prendas holgadas y neutras que se elevan de forma rápida con un saco.

PRINCIPIOS	PALETA
Priorizar la comodidad física usando ropa holgada	Negro
Elegir un look informal agregando un saco como tercera pieza	Beige
Mantener una base de color neutra y libre de gráficos grandes	Blanco

SILUETAS	EVITAR
Jeans anchos	Ropa entallada
Sacos rectos	Estampados grandes
Ropa holgada	

TU REGLA

"Si me pongo un saco encima de mis jeans anchos y tenis, ya la hice para la oficina."

Qué tan seguro estoy: **90%** - Me faltó saber: Qué tipo de playeras o tops prefiero usar debajo del saco y si uso otros colores además de los neutros.

[Guardar este núcleo](#)

A full run: the text someone wrote, and the core that came back. The line under the card states the model's confidence (**90%**) and what it never got to ask.

Núcleos que ya se generaron
Solo se muestra el núcleo. Lo que cada persona escribió sobre sí misma no se publica nunca.

Estilo casual relajado basado en jeans anchos y tenis, que se eleva rápidamente agregando un saco.

Negro - beige - blanco

"Si me pongo un saco encima de mis jeans y mis tenis, ya la hice para la oficina."

Estilo casual y cómodo estructurado con prendas holgadas neutras y un saco para elevar el look.

Negro - beige - blanco

"Para la oficina solo me pongo un saco encima de mis jeans y tenis, y ya la hice."

Ropa casual y común para el diario con opciones sencillas para arreglarse un poco más.

"No me gusta complicarme, prefiero lo normal para el diario pero que se vea bien."

Estilo femenino y muy colorido enfocado en faldas largas, vestidos florados y accesorios llamativos sin usar negro.

Rojos - Amarillos

"Antes muestra que sencilla: prefiero usar sandalias aunque camine mucho y jamás me veré de negro."

Estilo casual y cómodo de tonos neutros con prendas amplias que se eleva rápidamente para la oficina usando un saco.

Negro - beige - blanco

"Si traigo jeans y tenis, solo me pongo un saco encima y ya estoy lista para la oficina."

Comodidad estructurada usando jeans anchos y tenis con un saco para verse arreglada sin esfuerzo.

Negro - beige - blanco

"Si me pongo un saco encima de mis jeans y tenis, ya estoy lista para la oficina."

Look casual y holgado que se eleva de forma sencilla usando un saco estructurado.

Negro - beige - blanco

"Si me pongo un saco encima de mis jeans y tenis, ya estoy lista para la oficina."

The public list below the form. It shows the generated cores — and never the text anyone wrote about themselves.

5 · GitHub commits — 15 this week (minimum 5)

HASH	WHEN	MESSAGE
a4e9890	Sep 2, 20:16	docs: build discipline packet, before writing code
e02b657	Sep 2, 20:33	feat: <code>/core</code> — the style core, a public generative module
1f3a887	Sep 2, 20:39	docs: evidence for <code>/core</code> — prompt log, three runs, wireframe
6320e68	Sep 2, 20:40	feat: the landing page links to the style core
b2b01f8	Sep 2, 20:40	feat: "I don't know you yet" now offers a way into the core
7a8c1ec	Sep 2, 20:54	fix: retry when Gemini answers 429 instead of giving up
3a8bbef	Sep 2, 22:23	test: skip with a stated reason when Gemini runs out of quota, don't fail
21dd04d	Sep 3, 09:44	docs: image-generated mockup for <code>/core</code>
ffb8afe	Sep 3, 09:56	docs: demo video script
2dc22d0	Sep 3, 10:19	docs: video script in English too
b368c10	Sep 8, 19:23	feat: demo video page at <code>/demo/semana-1</code>
6bb7fb1	Sep 8, 19:45	fix: don't blame the person who wrote the text when it is the provider that failed
4baa73e	Sep 8, 21:30	docs: the pre-submission verification, including what failed
80ab136	Sep 9, 11:22	docs: this submission packet, with the source to rebuild it
c92ec6d	Sep 11, 15:21	feat: the demo video now carries my own voice

The list ends at the commit that shipped the narrated video; the commit that publishes this document into the repository comes after it.

Platform
Solutions
Resources
Open Source
Enterprise
Pricing

vestiamx-code / closetai
Public

Code
Issues
Pull requests
Actions
Projects
Security and quality
Insights

Commits

main

All users
All time

Commits on Sep 11, 2026

feat(semata 1): el video demo ya lleva la voz de Tamara

vestiamx-code and claude committed 32 minutes ago · 2 / 2

c92ec6d

Commits on Sep 9, 2026

docs(semata 1): el paquete de entrega, con su fuente para poder rehacerlo

vestiamx-code and claude committed 2 days ago · 2 / 2

80ab136

docs(semata 1): la verificación previa a entregar, con lo que falló

vestiamx-code and claude committed 3 days ago · 2 / 2

4baa73e

fix: no culpar a quien escribe cuando el que falla es el proveedor

vestiamx-code and claude committed 3 days ago · 2 / 2

6bb7fb1

feat(semata 1): página del video demo en /demo/semata-1

vestiamx-code and claude committed 3 days ago · 2 / 2

b368c10

Commits on Sep 3, 2026

docs(semata 1): guion del video también en inglés

vestiamx-code and claude committed last week · 2 / 2

2dc22d0

docs(semata 1): guion del video demo

The public repository — github.com/vestiamx-code/closetai. Every commit carries a green CI run.

6 · Vercel deployments — 9 this week (minimum 2)

All Production, all Ready. Counted from the Vercel deployment list, not from the number of pushes.

WHEN	ENVIRONMENT	STATUS	WHAT WENT OUT
Sep 2, 20:34	Production	● Ready	The <code>/core</code> page itself
Sep 2, 20:39	Production	● Ready	The week's evidence documents
Sep 2, 22:24	Production	● Ready	Entry points from the landing and the stylist, plus the 429 retry — a bug found by testing, fixed and redeployed
Sep 8, 19:22	Production	● Ready	The evidence commits that were still only local
Sep 8, 19:23	Production	● Ready	The demo video page
Sep 9, 11:22	Production	● Ready	This submission packet and the source to rebuild it
Sep 11, 15:21	Production	● Ready	The demo video with my narration
Sep 8, 19:45	Production	● Ready	The honest error message — a defect found while verifying, fixed and redeployed

ClosetAI · Week 1 · Generative Core Agent · Tamara Muñoz Delgadillo

9

WHEN	ENVIRONMENT	STATUS	WHAT WENT OUT
Sep 8, 21:31	Production	● Ready	The verification record

The third row is Gate 3 of the week in one line: a test found a real bug, the bug was fixed, and the fix was deployed.

7 · Supabase evidence

The migration, and the rows it actually holds — queried live from the production database, not from a local copy.

Migration 20260902000005_core_outputs.sql

```
create table if not exists public.core_outputs (  
  id          uuid primary key default gen_random_uuid(),  
  user_id     uuid references auth.users(id) on delete set null,  
  entrada     text not null check (char_length(entrada) between 40 and 1000),  
  nucleo      jsonb not null,  
  modelo      text,  
  version_prompt int,  
  created_at  timestampz not null default now()  
);  
  
alter table public.core_outputs enable row level security;  
  
create policy "núcleos visibles para todos"  
  on public.core_outputs for select using (true);  
  
revoke insert, update, delete on public.core_outputs from anon, authenticated;  
grant  select on public.core_outputs to anon, authenticated;  
grant  all    on public.core_outputs to service_role;
```

Two security decisions the assignment did not ask for. Writes are revoked from `anon`: without that, anyone could insert rows straight into the Supabase API, skipping the zod contract and the rate limit. And the public query **does not select the `entrada` column** — what someone writes about themselves is theirs. The list shows the core, not the confession.

Live contents — 10 rows

ROW	CREATED	MODEL	CONF.	PALETTE RETURNED
e840b638	Sep 3, 02:31	gemini-3.5-flash	90 %	negro · beige · blanco
4558e2a2	Sep 3, 02:33	gemini-3.5-flash	90 %	negro · beige · blanco
696d5b4f	Sep 3, 02:35	gemini-3.5-flash	90 %	negro · beige · blanco ← run 1
67db6e80	Sep 3, 02:36	gemini-3.5-flash	90 %	Rojo · Amarillo ← run 2
79fc4840	Sep 3, 02:36	gemini-3.5-flash	25 %	(empty) ← run 3
f5ff07fc	Sep 3, 02:43	gemini-3.5-flash	90 %	negro · beige · blanco
17d53b4d	Sep 3, 15:53	gemini-3.5-flash	90 %	negro · beige · blanco

ROW	CREATED	MODEL	CONF.	PALETTE RETURNED
d8c9b18a	Sep 9, 01:30	gemini-3.5-flash	90 %	negro · beige · blanco ← written by the test suite
39642c71	Sep 9, 01:37	gemini-3.5-flash	90 %	negro · beige · blanco ← written by the test suite
4cbf3e39	Sep 9, 01:40	gemini-3.5-flash	90 %	negro · beige · blanco ← written by the test suite

Every row records which model produced it and which version of the prompt, so a change in output quality can be traced to a change in the prompt.

The last three rows were written by the end-to-end suite during the verification described in §9 — which is itself the evidence that the test does what it claims: it saves for real, against the production database, and then checks that the row count went up.

8 · Prompt log — 6 prompts (minimum 5)

Each one records what I asked, what came back, and the judgment I had to apply on top of it. The judgment line is the part that matters: it is where I disagreed with the obvious answer.

1 · Plan before programming · a4e9890

"Here is the Week 1 assignment and its rubric. Tell me what we already have and what is missing, before writing anything."

Result: an inventory against the rubric. Five generative engines already existed — but all of them behind a login, and `/core` returned 404. The rubric has a hard cap: no live page, maximum 5/10.

My judgment: I decided the "generative module" would be Appendix A3 of my master document — how ClosetAI infers style — and not a generic text extractor. A generic extractor would have satisfied the letter of the assignment without serving the product.

2 · The packet, before the code · a4e9890

*"Write the Build Discipline Packet for this feature and **commit it before** touching any code, so the history proves the order."*

Result: `docs/SEMANA-1-PACKET.md` with problem, user, success, UX, scope cut, spec with checkable criteria, architecture, stack, DevOps and test plan.

My judgment: the scope cut was written with reasons, not as a wish list. Not wiring the core into the stylist yet is deliberate: first I have to prove the core that comes out is good.

3 · Migration and contract · e02b657

"Write the `core_outputs` migration and the zod contract. Do not invent columns."

Result: migration 005 and `styleCoreSchema`, with a cap on the length of every list.

My judgment: two security decisions the assignment never asked for. The table revokes `insert` from `anon` — otherwise anyone writes rows around the contract — and the public query **does not select the `entrada` column**: what someone writes about herself is hers.

4 · The module's prompt · e02b657

"Write the extraction prompt. Make it not invent: if the text is vague, it should say so."

Result: `src/lib/ai/prompts/core.ts`, versioned like the others. It explicitly forbids magazine filler ("timeless elegance", "versatile and sophisticated") and requires the model to declare `confianza` and `falta`.

My judgment: temperature 0.45, chosen by testing. At 0 the model repeated the same five phrases for different inputs; higher, it started inventing colors nobody had mentioned.

5 · Page and saving · e02b657

"Build `/core`: form, structured card, save button and the list of saved cores. Public, no session."

Result: the page placed outside the `(app)` route group so it does not inherit the session requirement, with the Server Actions `generarNucleo` and `guardarNucleo`.

My judgment: `guardarNucleo` **re-validates** whatever arrives from the browser. A Server Action can be called by direct POST; trusting that the payload is still the one we generated would leave the door open.

6 · Test against the real site · 1f3a887

"Run the three tests against production, not localhost, with different inputs — including one that is vague on purpose."

Result: all three passed. The third is the one that matters: with vague text the module returned 25% confidence, an empty palette, and said what it had failed to ask.

My judgment: two of my own tests were badly written before they were written well, and both failed the same way: they asserted text that **also exists on the freshly loaded page**. See §10.

9 · Test evidence — 3 live runs (minimum 3) + 13 automated tests

Against the live site, `closetai.lat/core` — not localhost. Each run left its row in `core_outputs` (§7), so these are checkable against the database.

#	INPUT	EXPECTED	ACTUAL
1	Rich, specific text: comfort, black and beige, wide jeans, sneakers, a blazer for the office	A complete core, high confidence, a palette containing the colors she named	6.4 s · 90 % · 3 principles · palette <code>negro · beige · blanco</code> · 3 silhouettes · 2 things to avoid. Nothing failed.
2	A deliberately opposite style: color, floral dresses, big earrings, hates black, sandals	A core different from the previous one. If it came out similar, the module would be returning templates	7.2 s · 90 % · 4 principles · palette <code>Rojo · Amarillo</code> · no black anywhere. Nothing failed.
3	Deliberately vague text: "I like to look good, nothing out of this world, just normal for every day"	That it admits it does not know , instead of filling in	7.5 s · 25 % · 2 principles · empty palette, empty silhouettes, empty avoid-list · states what it never got to ask. Nothing failed.

Run 3 is what was actually being tested. An extractor that returns a pretty palette to someone who never mentioned a color feels intelligent for ten seconds and false forever. The prompt forbids the filler and the result confirms it: with little information the lists come back empty and confidence drops to 25%.

Automated tests added this week

8 unit tests on the contract — that it rejects fewer than 2 principles, more than 5, a confidence value out of range, an array containing several cores, and that it distinguishes the model refusing from a formatting error.

5 end-to-end tests — that `/core` loads **with no session** (that is the feature, not a detail); that it rejects short text before spending a model call; that it generates the card with all its blocks and saves it to `core_outputs`; that the cost is logged; and that **the public list never shows what the person wrote**.

Full verification on submission day — against the live site

WHAT WAS RUN	RESULT
Unit tests (<code>vitest run</code>)	✓ 23 pass
End-to-end, one spec at a time	✓ all pass
End-to-end, all 58 in parallel	✓ 46 pass · 12 skipped for Gemini quota · 0 fail
<code>npm run build</code>	✓ compiles
<code>/core</code> with no session	✓ HTTP 200
<code>/demo/semana-1</code>	✓ HTTP 200

The 12 skips are the tests that depend on the reasoning model. Gemini's free tier allows 20 calls per minute and the suite exhausts it. **Skipped is not passed** — the report states each skip and its reason, and nobody collects a green they did not earn.

Being honest about how fragile this is: repeating the parallel run, between zero and four tests fail depending on the run — always ones that go through login or through the model. Run alone or in series they always pass, and the product verified by hand on the live site does not fail. It is the suite contending with production, not the code. It is written down rather than hidden.

10 · Iteration log — what changed after testing

Two badly written tests before one written well · Sep 2

Observed: the `/core` test failed while the card was visibly on screen. The module worked; the test did not.

First attempt. It asserted `getByText("Tu núcleo de estilo")`. It always failed, even with the card right there. The heading uses CSS `uppercase`, and Playwright reads the text **already transformed**: it was comparing "Tu núcleo de estilo" against "TU NÚCLEO DE ESTILO".

Second attempt. I made it case-insensitive. Then it passed — but it passed **instantly, in 0.1 s**, with no generation having happened at all. The page's intro paragraph says *"I distill your style core: your principles, your palette"*, so the assertion matched the freshly loaded page. **The test was green and was testing nothing.**

What stayed. The assertions are anchored to the card element, not to the page: `page.locator("article")`. Now there is no way for them to match the intro text.

Why it matters: this is exactly the same failure as Week 0 in a different disguise. There, the test asserted "the page loads without error" — and the login screen loads perfectly. Here it asserted text that already existed before anything

was generated. Both times the assertion could be satisfied **without the feature working**.

And for a while I believed the bug was in the product: I instrumented the Server Action step by step to find where it was hanging. The logs showed `model responded, ok = true` in 7.7 s — the module had been fine the whole time. **When a test and the product contradict each other, the test is a suspect too.**

A second thing testing changed · Sep 2

The runs hit Gemini's free-tier limit — 20 requests per minute — and came back `429 RESOURCE_EXHAUSTED`. I did not treat it as flakiness: I instrumented the call and read the actual limit off the response. Two changes came out of it, both deployed (§6): the call now retries, honoring the `retryDelay` the API itself asks for, and the end-to-end tests **skip with a stated reason** when quota is gone instead of reporting a failure that is not one.

Three red tests on submission day, and none of them was the code · Sep 8

Verifying the full suite against production before assembling this packet, three tests failed. All said the same thing: navigating to a private route bounced to `/entrar`.

The first move was not to believe the test. I opened the live site with a freshly created account, waited the way a person waits, and walked all six private routes. All fine. The product was not broken.

What was actually happening. I instrumented the browser and it came out clean: at the instant the URL changes to `/closet`, **the browser does not have the session cookie yet**. It arrives about 300 ms later — the redirect is done by the router using the payload in the same Server Action response, while the `Set-Cookie` headers are processed separately. With a 0 ms wait, all six routes failed; with 300 ms, none did. No person can click inside that gap. Playwright can.

What stayed. A single `entrarConSesionLista()` in `e2e/entorno.ts` that does not return until the cookie exists. Three specs each had their own copy of the login step — and all three failed for the same reason; now there is one. And the test now asserts something nobody was checking: **that logging in leaves a cookie**. I verified it by deliberately sabotaging the pattern it searches for, to watch it fail with the message I wrote rather than some other one.

The second finding. The Gemini quota gate in `core.spec.ts` sat in a `beforeAll` and skipped all five tests. Only one calls the model. The other four — that `/core` loads with no session, that it rejects short text, that the cost is logged, that the list does not expose what someone wrote — never touch Google. The central claim of the week was hiding behind a provider's quota.

The third is a product defect, and it is the one I care about most. When Gemini answered 429, the page told the person *"I couldn't read what you wrote properly."* Her text was fine; the one who did not answer was the model. There is now a separate reason, `unavailable`, and the page tells the truth: the model is overloaded, try again in a minute. Blaming someone for a failure that is not theirs costs more than the technical error.

11 · Demo video — 2:14

Hosted on the product's own domain rather than a third-party account that could expire or change its terms:
closetai.lat/demo/semana-1

The narration is my own voice. I recorded it in one take and read faster than the screen recording ran, so rather than let the two drift apart, the take was split at my own pauses and each block was placed on the moment it describes — "I save it" over the save button, "the repo is public" over GitHub. The only things cut from the screen recording were stretches where the picture was frozen.

ClosetAI

Semana 1 • El núcleo de estilo

Recorrido de 2:14 por `/core` en producción — sin cuenta, sin subir fotos y sin nada simulado.

0:00 / 2:14

Qué muestra

- El problema: la app pedía fotografiar diez prendas antes de dar nada
- `/core` abriéndose en el dominio propio, sin sesión iniciada
- Un texto escrito en palabras normales, no en vocabulario de moda
- El núcleo estructurado: esencia, principios, paleta, siluetas, qué evitar y una regla
- El guardado de la fila nueva en la tabla `core_output` de Supabase

The video page, live. The Week 0 video stays at `/demo`; each submission has its own stable link.

The walkthrough covers, in order: the problem the feature attacks; `/core` opening on the custom domain with no session; a text written in ordinary words; the structured core; the save and the new row in Supabase; the public list showing the core but never the input text; and the vague-input run where the module admits it does not know.

12 • Human Decision Note

The brief asked me to turn a prior methodology into a generative module. I could have built a generic text extractor and met the letter of it. I chose Appendix A3 of my master document — how ClosetAI infers someone's style — because it answers the strongest objection my value proposition had already received: my app makes you photograph ten garments before it gives you anything. `/core` gives you something useful in one minute, with no account and no photos.

The decision I am most sure of is the smallest one. The prompt forces the model to declare how confident it is and what it never got to ask, and it forbids magazine filler like "timeless elegance". I tested it with a deliberately vague text and it came back at 25% confidence, with the palette and the silhouettes empty. An extractor that invents a palette for someone who never mentioned a color feels clever for ten seconds and false forever.

I rejected requiring an account in order to save. That is the comfortable choice, and it is exactly the problem this page attacks.

One correction cost me. Twice I wrote tests that were green without testing anything. The second passed in a tenth of a second, because the text I was asserting already existed in the page's intro paragraph. It is last week's failure wearing a different disguise. I have learned to distrust a test that passes too fast.

239 words

Build gates

GATE	REQUIRED	EVIDENCE
1 — Plan	Build Discipline Packet complete before coding	§2 · commit <code>a4e9890</code> , 17 minutes ahead of the first feature commit
2 — Build	Code, commits and deploys completed	§5 · 15 commits · §6 · 9 production deployments
3 — Test	Tests documented, bug fixed, redeploy completed	§9 · 3 live runs · §10 · three defects found by testing, fixed and redeployed (Sep 2, 22:24 and Sep 8, 19:45)
4 — Explain	Demo and human decision note submitted	§11 · 2:14 video · §12 · 239 words