

Is the problem real?

With sources.

A public research module for ClosetAI where every data point links to the page that proves it — and where the AI can only reason over data a person already verified.

Tamara Muñoz Delgadillo · Universidad Iberoamericana
Submitted September 25, 2026 · Product live at **closetai.lat**

LIVE PAGE	closetai.lat/research
GITHUB	github.com/vestiamx-code/closetai
DEMO VIDEO	closetai.lat/demo/semana-2

THE WEEK IN ONE LINE. Before building on my market thesis, I checked it source by source. Nine claims in my own master document did not survive. The biggest one: my thesis was "there is no closet app in Spanish" — **there are several**. The gap is not the language; it is Mexico: sizes, stores and prices in pesos. I also found an app that is already called "Closet AI". Building three more weeks on an unverified thesis would have been the most expensive mistake this product could make.

1 • The rubric, category by category

CATEGORY	PTS	EVIDENCE	\$
Build discipline before coding	1.5	Packet committed at 11:37 (<code>e3f65d3</code>); first line of code at 11:51 (<code>1d0f355</code>). Public history.	2
UX planning and image-generated mockup	1.0	Image-generated mockup (Gemini) · wireframe with implementation notes and scope cut	3
Product spec and acceptance criteria	1.0	11 requirements, each with a checkable criterion	2
Architecture and tech stack clarity	1.0	Data flow from browser to Postgres, with the contract gate marked; stack table with the why	2
Working deployed product	2.0	<code>closetai.lat/research</code> loads with no session; 14 of 14 end-to-end tests pass against production	4, 10
Coding/build evidence	1.0	15 commits (min. 5) · 14 deployments (min. 2) · 10 coding-agent prompts	6–9
Testing and iteration	1.0	50 unit tests · end-to-end against production · 3 live runs · 5 defects found by testing, fixed and redeployed · 1 real human conversation , in person, with consent	10–12
Human judgment and explanation	1.0	9 corrections to my own research · decision note, 240 words	5, 14
Demo clarity	0.5	2:30 walkthrough recorded against production, narrated by me	13

Required features and coding tasks from the week overview

REQUIRED	WHERE IT IS
Research intake	A research question (20–500 characters) that returns a report citing only verified sources
5 global examples	Benchmark cards: Whering, Stylebook, Style DNA, Indyx, Google try-on
Mexico localization	Mexico section: 77 M digital buyers, 59% buy clothes online, 35% "very satisfied", GoTrendier, Spanish-language apps
8 competitors/substitutes	11 in a table: 9 competitors and 2 substitutes
Filter/search table	Filter by type; search that ignores accents ("mexico" finds "México")
Risk map	7 risks on a 3×3 probability × impact grid, each linked to its evidence
Saved research output	<code>research_outputs</code> table; saved reports listed on the page
1 human validation conversation	Done in person on Sep 21, with consent; published on the page without a name (§11)
Create competitor table · add filter/search · add benchmark cards · save research record · add dashboard widget	All live. The dashboard widget is the summary panel at the top, computed from the database

2 · Build Discipline Packet

Written before touching code. Committed at 11:37 (`e3f65d3`); the first file of the feature arrived at 11:51 (`1d0f355`). The research it stands on was verified first, in the same commit.

Problem

ClosetAI was designed on a market thesis — "nobody serves Mexico" — written in August from desk research. When I checked it, part of it was false: **there are closet apps in Spanish**, and **there is already an app called "Closet AI"**. Building three more weeks on an unverified thesis is the most expensive mistake a new product can make. And nobody outside the project — a professor, an investor, a potential user — could check today why ClosetAI would have a place in the market.

User

(1) **Me, the founder**, deciding with evidence what to build, at what price and under what name. (2) **Whoever evaluates ClosetAI** — professor, investor — who needs to see the proof, not believe it. (3) **The person interviewed**, whose real conversation becomes evidence, with consent.

Success — what has to work by the end of the week

- `/research` loads in production with no session.
- Every data point links to its source and carries a verification date.
- 5 global examples, Mexico localization, and at least 8 competitors and substitutes in a table that can be filtered and searched.
- A risk map where each risk points to the evidence behind it.
- A research question produces a report that **only uses verified data** and cites every claim; it can be saved and lands in Supabase.
- **One real conversation with a person**, recorded with consent. Until it exists, the page says "pending": it is never simulated.

UX concept

One page, top to bottom, ordered the way an argument reads: summary panel → research question and report → 5 global examples → Mexico → competitors and substitutes → risk map → real conversation → saved reports → what we discarded. Wireframe in §3.

Scope cut — what looked nice and is deliberately not built

- **An AI that researches the internet.** A model that searches and summarizes invents figures with total confidence — and in this same research, the search engine's summaries got two figures wrong. The AI only reasons over data a person verified.
- **Automatic competitor updates** (scraping): fragile, against several stores' terms, and it changes data nobody reviews.
- **Editing the research from the page:** data enters through a migration, with its source.
- **A charting library:** a CSS grid is enough for 7 risks.

- **More than one validation conversation:** the assignment asks for one; doing it well beats doing five in a rush.

Product spec and acceptance criteria

#	REQUIREMENT	ACCEPTANCE CRITERION (CHECKABLE)
1	Public page	GET /research returns 200 with no session cookie
2	Summary panel	Counts come from the database, not from the code
3	5 global examples	5 cards, each with key fact, lesson, source link and date
4	Mexico localization	A Mexico section, each data point with its source
5	Competitors and substitutes	A table with at least 8 rows of type competitor or substitute
6	Filter and search	Filtering by type shows only that type; "mexico" finds "México"; no match shows a message, not an empty table
7	Risk map	Each risk placed by probability × impact, with its mitigation and evidence
8	Research question	20–500 characters; the report passes a zod contract; every cited source exists in the data ; a claim with no source is rejected; if the data is not enough, the report says so
9	Save	Writes a row to <code>research_outputs</code> and appears in the list without reloading
10	Human validation	Shows the real, consented conversation; if none, says "pending"
11	Cost and limit	Every report logs its cost; one IP cannot generate more than 10 per hour

Architecture

```

Browser (/research, no session)
  |
  ▼
Server Component – reads research_sources · research_risks
                    validation_conversations · research_outputs
  ├── Table (client): filter and accent-insensitive search over the loaded rows
  ├── Risk map: 3×3 CSS grid
  └── Form → Server Action generarInvestigacion()
      ├── validates the question with zod · checks the per-IP limit
      ├── reads the verified data from Supabase
      ├── calls Gemini with RESEARCH_PROMPT + that data, identified by id
      │   └── if the model is saturated or times out (25 s) → lighter backup model
      ├── contract: zod schema + every cited id must exist ← if not, nothing is shown
      ├── logs the cost in api_costs
      └── ▼ (the person decides to save)
          guardarInvestigacion() ← re-checks the citations: reachable by direct POST
              └── inserts into research_outputs

```

Why the AI cannot search the internet: the report can only cite ids that exist in the table. An unverified data point cannot appear, because the contract rejects the whole report.

Tech stack

TOOL	FOR WHAT	WHY THIS ONE
Next.js 16 · Server Components and Actions	Page and logic	The project's stack; the Gemini key never leaves the server
Supabase Postgres	4 new tables	This week the research is evidence: it has to live in the database, not in the code
Gemini <code>flash</code> , with <code>flash-lite</code> as backup	Report from the data	Reasoning over short text; the backup was added after measuring the main model at 86 s
zod	Report and citation contract	The fence between what the model says and what is shown
Tailwind v4	Styles and risk map	No chart library needed. No new dependencies

⚙️ DevOps

- Public repo `vestiamx-code/closetai`, branch `main`, automatic deploy to Vercel.
- Migration `006_research` creates the 4 tables and seeds the verified data with source and date; applied to production **before** deploying, and verified against the database.
- RLS on all 4 tables: public read, server-only write. No new environment variables.

✍️ Test plan

Automated: unit tests for the contract (it rejects an invented source and a claim with no source) and for accent-insensitive search; end-to-end tests that the page loads with no session and at least 8 rows, that a search with no match shows the message, that the risk map shows every risk, and that generating and saving a report writes a row. **Three live runs** against production, including **one question the data cannot answer**. **One real conversation** with a person in the segment, asking about past behavior rather than opinions of the idea.

🤖 Coding agent prompt

Build the public page /research in ClosetAI (Next.js 16, App Router, Server Components and Actions, Tailwind v4, Supabase).

Source of truth: docs/evidencia/INVESTIGACION-SEMANA-2.md. Do not add any data point that is not there, with its source.

First the migration 006_research: tables research_sources, research_risks, validation_conversations and research_outputs, with public-read / server-write RLS. Seed the verified data with its source URL and date. The validation table requires consent and is NOT seeded: Tamara fills it with a real conversation.

The page shows, in order: a summary panel with counts computed from the database; a research question form; 5 global benchmark cards; the Mexico section; the competitor and substitute table with type filter and accent-insensitive search; a 3x3 risk map; the validation conversation or "pending"; saved reports; and what was discarded as unverifiable.

The report is generated by Gemini ONLY from the database rows, identified by id. Zod contract: every claim carries at least one source and every source must exist in the data; if the data is not enough, the report says so. Version the prompt in src/lib/ai/prompts/research.ts. Log the cost in api_costs and limit to 10 reports per IP per hour, as in /core.

3 · UX: wireframe and mockup

The wireframe fixes the order of the argument and carries the implementation notes and scope cuts next to each section.

Wireframe · /research

ClosetAI · Semana 2 · Una sola página, ordenada como se lee un argumento: ¿el problema es real? → ¿quién más? → ¿qué nos puede matar? → ¿qué dice una persona?

¿EL PROBLEMA ES REAL? CON FUENTES.

1 PANEL DE RESUMEN

9 Competidores

2 Sustitutos

5 Referentes

7 Riesgos

Riesgo más alto ·

Validación · Pendiente / 1 conversación

2 PREGUNTA DE INVESTIGACIÓN

ejemplo 1

ejemplo 2

ejemplo 3

Investigar

INFORME · SÍNTESIS DE IA A PARTIR DE FUENTES CITADAS

· hallazgo

Stylebook

· hallazgo

Style DNA

Hueco

Riesgo

Falta validar

Guardar este informe

3 5 REFERENTES GLOBALES

Whering

fuente ↗ · fecha

Stylebook

Style DNA

Indyx

Google

4 MÉXICO

77 M compradores

59% compra ropa

35% muy satisfecha

GoTrendier · 8 M

5 COMPETIDORES Y SUSTITUTOS

Todos

Competidores

Sustitutos

Buscar...

Nombre	Tipo	Dato verificado + fuente	Precio	Amenaza
Whering	competidor		Freemium	●●●
Acloset	competidor		\$3.99-24.99	●●○
GoTrendier	sustituto		14% + \$14	●○○

... 8 o más filas · sin resultados: aviso en lugar de tabla vacía

6 MAPA DE RIESGOS

Impacto alto

monetizar

fricción

Impacto medio

satisfacción

Google | nombre

Impacto bajo

Prob. baja

media

alta

7 VALIDACIÓN CON UNA PERSONA REAL

"cita textual de la persona"

"otra cita"

Confirmó

Contradijo

Sorprendió

Si no hay conversación: "Pendiente" — nunca se simula

8 GUARDADOS

9 LO QUE DESCARTAMOS

NOTAS DE IMPLEMENTACIÓN

1 · Panel

Todos los conteos salen de Supabase. Ninguno está escrito en el código.

2 · Informe

Gemini recibe solo las filas verificadas, con su id. El contrato zod rechaza el informe si cita un id que no existe. Si los datos no alcanzan, lo dice.

3 · Referentes

Filas con referente_global = true. Cada tarjeta: dato, lección, fuente y fecha.

4 · México

Filas con de_mexico = true. Traducir no es localizar.

5 · Tabla

Componente de cliente: el filtro y la búsqueda corren sobre las filas ya cargadas. La búsqueda ignora acentos.

6 · Mapa

Cuadrícula de CSS, sin librería. Debajo, la lista con mitigación y evidencia de cada riesgo.

7 · Validación

Tabla con check de consentimiento. No se siembra: la llena una conversación real.

RECORTES DE ALCANCE

Sin IA que busque en internet

El buscador se equivocó dos veces en esta investigación.

Sin scraping ni edición pública

Los datos entran por migración, con su fuente.

Sin librería de gráficas

Una cuadrícula alcanza para 7 riesgos.

Wireframe of /research, with implementation notes and scope cuts.

ClosetAI · Week 2 · Tamara Muñoz Delgadillo · closetai.lat

Page 7 of 22




Image-generated mockup (Gemini), desktop and phone. A concept, not data: every number, source and even the “real conversation” quote in this image were **invented by the image model** as placeholder text — none of them are in the product. The verified figures are in §4 and §5, and the real conversation is in §11. It is, in its own way, the argument of the week: left alone, a generative model fills gaps with confident, made-up sources. That is why the AI in `/research` can only cite rows a person verified.

4 • The live page

```
$ curl -s -o /dev/null -w "HTTP %{http_code}\n" https://closetai.lat/research
HTTP 200
```

ClosetAI

INVESTIGACIÓN

Probar el núcleo de estilo

¿El problema es real? *Con fuentes.*

Quién más resuelve lo mismo, qué usa la gente en su lugar, qué pasa en México y qué podría matar a ClosetAI. Cada dato enlaza a su fuente; lo que no pudimos verificar está al final, fuera de la tabla.

Verificado el 14 de septiembre de 2026 · [cómo se verificó cada dato](#)

9

Competidores

2

Sustitutos

5

Referentes globales

7

Riesgos mapeados

RIESGO MÁS ALTO

La gente abandona antes de catalogar su ropa

VALIDACIÓN CON PERSONAS

1 conversación real · 9 informes guardados

Hazle una pregunta a la investigación

Pregunta de investigación

La respuesta sale solo de los datos verificados de esta página. La pregunta y el informe se publican si los guardas: no escribas datos personales.

¿Hay competidores hechos para México?

Top of the page: the question it answers, the verification date, and the summary panel. Every number in the panel is computed from the database.

5 referentes globales

Lo que ya funcionó —o no— en otros mercados, y lo que le enseña a ClosetAI.

COMPETIDOR · REINO UNIDO

Whering
10 millones de usuarios, sobre todo Gen Z; ronda de 7 M USD en julio de 2026 con eBay Ventures y Google AI Futures Fund
Freemium

Lección: El clóset gratis escala: es la puerta de entrada, no el producto
[tech.eu](#) ✓ · verificado el 2026-09-14

COMPETIDOR · ESTADOS UNIDOS

Stylebook
Pago único, sin suscripción; solo funciona en iPhone y iPad
\$4.99 USD una sola vez

Lección: El pago único barato es un modelo que existe y dura
[Stylebook \(FAQ\)](#) ✓ · verificado el 2026-09-14

COMPETIDOR · REINO UNIDO

Style DNA
3 M USD de ingreso anual a fines de 2023; 300 mil usuarios activos y 70 mil de pago
Freemium más comisión

Lección: El nicho monetiza, pero no sin capital: ahorros, luego ángeles y un fondo
[Starter Story](#) ✓ · verificado el 2026-09-14

COMPETIDOR · ESTADOS UNIDOS

Indyx
Catalogar tu clóset es gratis; por 295 USD ellos catalogan 100 prendas por ti
Estilista desde \$15 USD al mes

Lección: Hay quien paga 295 USD por no fotografiar su ropa: la fricción es real
[Indyx](#) ✓ · verificado el 2026-09-14

SUSTITUTO · ESTADOS UNIDOS

Try-on de Google
Google cerró Doppl el 30 de abril de 2026 y pasó el try-on a su buscador; reportado en Latinoamérica en mayo de 2026, sin confirmar México por nombre
Gratis

Lección: Si Google regala el try-on, el valor tiene que venir de tu clóset
[Google Labs](#) ✓ · [otra fuente](#) ✓ · verificado el 2026-09-14

5 global examples: verified key fact, lesson for ClosetAI, source and date.

México

Traducir no es localizar. Lo que dicen los datos sobre el mercado donde ClosetAI sí puede ser el primero.

Apps de clóset en español
Existen en App Store con interfaz en español; ninguna de las revisadas se presenta con tallas, tiendas o precios de México

Traducir no es localizar: el hueco no es el idioma, es México
[App Store](#) ✓ · [otra fuente](#) ✓ · verificado el 2026-09-14

GoTrendier
Más de 8 millones de usuarios; 4.6 estrellas con 74 mil calificaciones en App Store México

La segunda mano ya tiene escala en México: integrarla, no competirle
[App Store México](#) ✓ · verificado el 2026-09-14

Compradores digitales en México
Más de 77 millones en 2025; eran 37 millones en 2018

El canal existe y se duplicó en siete años
[Milenio con datos de AMVO](#) ✓ · verificado el 2026-09-14

Ropa en línea en México
El 59% de los compradores digitales compra ropa

La ropa ya se compra en línea: no hay que convencer a nadie de eso
[Milenio con datos de AMVO](#) ✓ · verificado el 2026-09-14

Experiencia de compra de moda en línea
Más del 50% planeó comprar moda en línea, pero solo el 35% calificó su experiencia como muy satisfactoria

Hay interés y hay insatisfacción: ese espacio es el hueco
[AMVO](#) ✓ · verificado el 2026-09-14

Mexico localization.

ClosetAI • Week 2 • Tamara Muñoz Delgadillo • closetai.lat

Page 9 of 22

Competidores y sustitutos

Competidor: resuelve lo mismo. Sustituto: resuelve la necesidad de otra forma.

TodosCompetidoresSustitutos

Buscar: nombre, país, precio...

11 de 11

NOMBRE	TIPO	DATO VERIFICADO	PRECIO	AMENAZA
Whering Reino Unido	Competidor	10 millones de usuarios, sobre todo Gen Z; ronda de 7 M USD en julio de 2026 con eBay Ventures y Google AI Futures Fund tech.eu ↗	Freemium	●●● Alta
Acloset Corea del Sur	Competidor	4.4 estrellas con 5.1 mil calificaciones en App Store; la ficha declara que tiene publicidad App Store ↗	\$3.99 a \$24.99 USD al mes	●●● Media
Stylebook Estados Unidos	Competidor	Pago único, sin suscripción; solo funciona en iPhone y iPad Stylebook (FAQ) ↗	\$4.99 USD una sola vez	●●● Baja
Style DNA Reino Unido	Competidor	3 M USD de ingreso anual a fines de 2023; 300 mil usuarios activos y 70 mil de pago Starter Story ↗	Freemium más comisión	●●● Media
Alta Estados Unidos	Competidor	Ronda semilla de 11 M USD en junio de 2025, liderada por Menlo Ventures y con el fondo de la familia Arnault TechCrunch ↗	—	●●● Media
Doji Estados Unidos	Competidor	Ronda semilla de 14 M USD en mayo de 2025, liderada por Thrive Capital; por invitación en más de 80 países TechCrunch ↗	—	●●● Media
Indyx Estados Unidos	Competidor	Catalogar tu clóset es gratis; por 295 USD ellos catalogan 100 prendas por ti Indyx ↗	Estilista desde \$15 USD al mes	●●● Baja
Apps de clóset en español Varios	Competidor	Existen en App Store con interfaz en español; ninguna de las revisadas se presenta con tallas, tiendas o precios de México App Store ↗	Variable	●●● Media
Closet AI – Outfit Planner	Competidor	Una app con el mismo nombre en iPhone; menos de mil descargas Ficha de la app ↗	Pro a \$4.99 USD al mes	●●● Media
Try-on de Google Estados Unidos	Sustituto	Google cerró Doppl el 30 de abril de 2026 y pasó el try-on a su buscador; reportado en Latinoamérica en mayo de 2026, sin confirmar México por nombre Google Labs ↗	Gratis	●●● Alta
GoTrendier México	Sustituto	Más de 8 millones de usuarios; 4.6 estrellas con 74 mil calificaciones en App Store México App Store México ↗	Comisión de 14% más \$14 MXN	●●● Baja

Competitors and substitutes: 11 rows, each with its source link.

Competidores y sustitutos

Competidor: resuelve lo mismo. Sustituto: resuelve la necesidad de otra forma.


TodosCompetidoresSustitutos

mexico


3 de 11

NOMBRE	TIPO	DATO VERIFICADO	PRECIO	AMENAZA
Apps de clóset en español Varios	Competidor	Existen en App Store con interfaz en español; ninguna de las revisadas se presenta con tallas, tiendas o precios de México App Store ↗	Variable	●●● Media
Try-on de Google Estados Unidos	Sustituto	Google cerró Doppl el 30 de abril de 2026 y pasó el try-on a su buscador; reportado en Latinoamérica en mayo de 2026, sin confirmar México por nombre Google Labs ↗	Gratis	●●● Alta
GoTrendier México	Sustituto	Más de 8 millones de usuarios; 4.6 estrellas con 74 mil calificaciones en App Store México App Store México ↗	Comisión de 14% más \$14 MXN	●●● Baja

Typing "mexico", with no accent, finds México.



The risk map: 7 risks by probability × impact; below, each one's mitigation and evidence.



On a phone, with no horizontal overflow.

Lo que descartamos

Pueden ser ciertos, pero no los pudimos comprobar en su fuente. Por eso no están arriba.

- Que Chicisimo haya sido la app de moda número uno en México en 2014.
- Que el 90% de las usuarias de GoTrendier sean mujeres.
- Que catalogar en Stylebook tome de 6 a 8 horas por cada 100 prendas.
- Que Doji haya cambiado la raza de usuarios en sus avatares.
- Las rondas de Daydream (50 M USD) y Phia (35 M USD).
- El ticket promedio de moda en línea en México de 1,300 MXN.

What we discarded: possibly true, but not verifiable at the source — so it is not in the data.

5 · What verifying the research changed

Every data point was checked by opening its source, not by reading a search engine's summary. That rule turned out to be necessary: the search engine summarized two figures wrong (Indyx's stylist price: it said \$25/month, the page says \$15; GoTrendier's commission: it said 20% + \$9, the developer says 14% + \$14).

#	MY MASTER DOCUMENT SAID	THE SOURCE SAYS
1	"There is no native Spanish/LATAM player"	There are closet apps in Spanish on the App Store. The thesis changes: the gap is Mexico, not the language
2	Style DNA is bootstrapped	Its founder: personal savings, then angel investors and a fund
3	Acloset: "€130/year paywall"	App Store: plans from \$3.99 to \$24.99 USD a month, up to \$147.99 a year
4	Alta and Doji raised in 2026	Both rounds were in 2025
5	Google's try-on "available in Mexico"	Reported for Latin America (May 2026); Mexico is not named
6	—	An app called "Closet AI" already exists — a new brand risk
7	Average online fashion ticket ~\$1,300 MXN	No source found — excluded
8	—	GoTrendier, Mexican, has over 8 million users: secondhand already has scale in Mexico
9	—	Indyx charges \$295 USD to catalog 100 items for you: cataloging friction is real

Discarded as unverifiable (they may be true; they are not in the data): that Chicisimo was the #1 fashion app in Mexico in 2014; that 90% of GoTrendier users are women; that cataloging 100 items in Stylebook takes 6–8 hours; that Doji changed users' race in avatars; the Daydream and Phia rounds; the \$1,300 MXN average ticket.

6 · GitHub commits — 15 this week (minimum 5)

HASH	WHEN	MESSAGE
e3f65d3	Sep 14, 11:37	docs: Build Discipline Packet and verified research, before writing code
1d0f355	Sep 14, 11:51	feat: <code>/research</code> — research and benchmarking with verified sources
1b031c3	Sep 14, 11:53	docs: <code>/research</code> wireframe and the week's prompt log
bb9ac53	Sep 14, 12:16	fix: a Gemini 503 no longer asks the person to rephrase their question
23abff3	Sep 14, 12:33	fix: if Gemini is saturated, the lighter model answers instead of making people wait
4854dd1	Sep 14, 12:41	test: the quota gate checks both models, not just the main one
46a1e64	Sep 14, 12:43	fix: prompt v2 — the report does not generalize beyond the data
f9ffc84	Sep 14, 12:55	docs: test and iteration evidence, and the per-IP limit in e2e
11f1e28	Sep 18, 17:04	fix(db): every new account was born with a mis-encoded city

HASH	WHEN	MESSAGE
7494e17	Sep 21, 11:47	docs: the real validation conversation and the decision note
e857c80	Sep 21, 11:56	perf(ai): the backup model steps in at 15 s, not 25
a2badee	Sep 21, 12:02	feat: demo video page at <code>/demo/semana-2</code>
7eeca68	Sep 21, 12:05	docs: decision note approved by me
539ad85	Sep 21, 12:13	docs: <code>/research</code> mockup generated with Gemini
87c2711	Sep 25, 16:47	feat: the demo video now carries my own narration

Platform Solutions Resources Open Source Enterprise Pricing

Search Sign in Sign up

vestiamx-code / closetai Public

Notifications Fork 0 Star 0

Code Issues Pull requests Actions Projects Security and quality Insights

Commits

main

All users All time

Commits on Sep 25, 2026

feat(semana 2): el video demo ya lleva la voz de Tamara

vestiamx-code and claude committed 4 minutes ago · 2 / 2

87c2711

Commits on Sep 21, 2026

docs(semana 2): mockup de /research generado con Gemini

vestiamx-code and claude committed 4 days ago · 2 / 2

539ad85

docs(semana 2): nota de decisión aprobada por Tamara

vestiamx-code and claude committed 4 days ago · 2 / 2

7eeca68

feat(semana 2): página del video demo en /demo/semana-2

vestiamx-code and claude committed 4 days ago · 2 / 2

a2badee

perf(ia): el respaldo entra a los 15 s, no a los 25

vestiamx-code and claude committed 4 days ago · 2 / 2

e857c80

docs(semana 2): la conversación real de validación y la nota de decisión

vestiamx-code and claude committed 4 days ago · 2 / 2

7494e17

Commits on Sep 18, 2026

fix(db): toda cuenta nueva nacía con la ciudad mal codificada

The public repository. Every commit carries a green CI run.

7 · Vercel deployments — 14 this week (minimum 2)

All Production, all Ready, with exact times read from Vercel.

WHEN	STATUS	WHAT WENT OUT
Sep 14, 11:37	● Ready	The packet and the verified research
Sep 14, 12:08	● Ready	<code>/research</code> — after its migration was applied and verified

WHEN	STATUS	WHAT WENT OUT
Sep 14, 12:16	● Ready	Fix: 503 recognized as a provider failure
Sep 14, 12:33	● Ready	Fix: backup model with a 25 s limit per attempt
Sep 14, 12:41	● Ready	Test fix: the quota gate
Sep 14, 12:44	● Ready	Prompt v2
Sep 14, 12:55	● Ready	Evidence, and the per-IP limit in tests
Sep 18, 17:04	● Ready	Fix: the mis-encoded city for new accounts
Sep 21, 11:48	● Ready	The validation conversation and the decision note
Sep 21, 11:56	● Ready	The backup model at 15 s
Sep 21, 12:02	● Ready	The demo video page
Sep 21, 12:06	● Ready	The approved decision note
Sep 21, 12:14	● Ready	The Gemini mockup in the evidence
Sep 25, 16:47	● Ready	The demo video with my narration

8 · Supabase evidence

Migration `20260914000006_research`: 4 tables, RLS, and the verified data seeded with its source URL and verification date. Checked against the database, not the screen.

WHAT WAS CHECKED	RESULT
Rows in <code>research_sources</code> / <code>research_risks</code>	15 / 7
<code>validation_conversations</code> , <code>research_outputs</code> created	✓
Public read (anon key)	✓ allowed
Public write (anon key)	✓ rejected — 401 permission denied
A conversation with <code>consentimiento = false</code>	✓ rejected by the database itself — 23514
10 reports per IP per hour	✓ the 11th was blocked; <code>api_costs</code> logged exactly 10
Saved reports in <code>research_outputs</code>	7, each recording the model and prompt version

```
create table public.validation_conversations (  
  id uuid primary key default gen_random_uuid(),  
  fecha date not null,  
  perfil text not null,  
  citas text[] not null check (cardinality(citas) >= 1),  
  confirmo text, contradijo text, sorpresa text,  
  consentimiento boolean not null check (consentimiento), -- without consent, the row does not go in  
  created_at timestamptz not null default now()  
);
```

The research is evidence, so it lives in the database. Postgres cannot put a foreign key on an array, so an end-to-end test guarantees that every risk's evidence id exists.

9 · Prompt log — 10 prompts (minimum 5)

What I asked, what came back, and the judgment I applied on top.

1 · Verify before planning · `e3f65d3`

"Before building, check what research already exists in the master document and verify every data point at its source."

My judgment: open the source, never the search summary. It changed the thesis: the gap is Mexico, not the language.

2 · The packet, before the code · `e3f65d3`

"Write the Build Discipline Packet on the verified research and commit it before touching code."

My judgment: the most important scope cut was not letting the AI search the internet.

3 • Migration and seeded data • 1d0f355

"Seed only what is in the research document, with URL and date. The validation table is not seeded."

My judgment: consent is enforced by the database, not only by the interface.

4 • A contract that forbids invented sources • 1d0f355

"If the report cites an id that does not exist, reject it entirely."

My judgment: the prompt asks; the contract enforces. Temperature 0.2: the same evidence should give the same answer.

5 • The page • 1d0f355

"Build /research: panel from the database, question, 5 examples, Mexico, table, risk map, validation, discarded."

My judgment: search ignores accents; validation says "Pending" instead of faking it.

6 • Watch the tests fail before trusting them • 1d0f355

"Sabotage the code on purpose and check that the important tests fail for the right reason."

My judgment: a test you have never seen fail proves nothing.

7 • Reproduce the failure outside the browser • bb9ac53

"Don't change anything yet. Reproduce the call with the same question and data, and show me the raw response."

My judgment: diagnose before fixing. The raw response showed a 503 — the question was fine.

8 • Measure before adding a backup model • 23abff3

"Does the SDK allow a timeout? Is the lighter model answering while the main one isn't?"

My judgment: fall back only on provider failures; a bad request must not be hidden by switching models.

9 • Prompt v2 from a real run • 46a1e64

"The report said 'there are no local competitors' when the source says 'none of the ones reviewed'. Add the rule and repeat the same question."

My judgment: no contract catches a sentence that over-generalizes. A person has to read it.

10 • Read the screen of the failure, not just the message • 11f1e28

"Run the full suite against production. If something fails, don't call it flaky until you read what was on the screen."

My judgment: the bug that mattered was not the one that failed the test, but one sitting in its screenshot.

10 · Test evidence

Three live runs against closetai.lat/research (minimum 3)

#	QUESTION	EXPECTED	ACTUAL
1	"Are there competitors built for Mexico, or only translated apps?"	Cite the Mexico sources and keep their limits	2.9 s · cites 2 sources · over-generalized : "there are no local competitors", when the source says "none of the ones reviewed" → prompt v2
1 (v2)	The same	No longer claims "there are none"	4.9 s · "The Spanish-language closet apps reviewed do not present Mexican sizes, stores or prices" ✓
2	"How big is online clothing shopping in Mexico, and how satisfied are people?"	Exact figures, not rounded	3.3 s · cites the 3 AMVO data points · 77 M, 37 M (2018), 59%, over 50% and 35%, copied exactly ✓
3	"How many people in Guadalajara use apps to organize their clothes?" — outside the data, on purpose	Admit it does not know	2.3 s · "The verified data is not enough to answer this" · 0 citations, 0 invented figures · suggests a local survey ✓

INFORME

¿Hay competidores hechos para México, o solo apps traducidas?

Las apps de clóset en español tienen interfaz en nuestro idioma, pero ninguna de las revisadas se presenta con tallas, tiendas o precios de México. En contraste, GoTrendier sí opera en el país con más de 8 millones de usuarios. Por lo tanto, no hay competidores locales dedicados a organizar clósets, sino aplicaciones traducidas o plataformas de segunda mano.

Los datos verificados alcanzan para responder. Síntesis de IA a partir de las fuentes citadas.

Ninguna de las apps de clóset en español revisadas se presenta con tallas, tiendas o precios de México.

Apps de clóset en español ✓

GoTrendier tiene más de 8 millones de usuarios en México y 4.6 estrellas con 74 mil calificaciones en App Store México.

GoTrendier ✓

HUECO PARA CLOSETAI

Crear una app de clóset nativa para México que incluya tallas, tiendas y precios locales.

RIESGO

Que los competidores internacionales traduzcan su interfaz sin ofrecer un verdadero producto adaptado a México.

FALTA VALIDAR

Falta validar si existen startups locales menores que operen exclusivamente en México y cómo validarían su tracción; se investigaría directamente en tiendas de aplicaciones locales y registros de empresas.

Guardar este informe

INFORME

¿Cuántas personas en Guadalajara usan apps para organizar su ropa?

Los datos disponibles no incluyen cifras sobre cuántas personas en Guadalajara usan aplicaciones para organizar su ropa. La información sobre mercado solo abarca el total de compradores digitales y su comportamiento en México.

Los datos verificados no alcanzan para responder esto. Síntesis de IA a partir de las fuentes citadas.

FALTA VALIDAR

No hay datos sobre Guadalajara; se tendría que realizar una encuesta local o un estudio de mercado específico para esa ciudad.

Guardar este informe

Run 3: outside the data — no citations, no invented numbers.

Run 1 (v1): note "Falta validar" admits what the answer then over-states.

Automated tests

50 unit tests (27 new this week): the report contract — including **rejecting the whole report if it cites a source that does not exist** —, accent-insensitive search, risk ranking, and provider-failure handling with the literal 503 text from production. The key ones were verified by **sabotaging the code to watch them fail** for the right reason.

End-to-end, against production, on phone and desktop: `/research` **14 of 14**. Full suite on Sep 18: **67 pass**; 1 intermittent stylist test (4 of 4 when run alone); 4 skipped by design or in chain. That run also uncovered a real bug — see §12.

Re-verified on September 25, before submitting: 50 unit tests pass, and `/research` is 14 of 14 against production. In the first of three runs that afternoon, the risk-map test failed once on the phone profile and passed on every repeat,

alone and in the full file. The cause is not identified, so it is written here instead of being left out: an intermittent test is a test I do not trust yet.

11 • Human validation conversation

In person, on September 21, with consent to use it without a name — in this document and on the public page. Two limits, stated on the page itself: she is my mother, and at 57 she is outside my core 18–35 segment.

WHAT SHE SAID OR DID	
The problem	“Entre tanta ropa que tengo no sé qué ponerme, entonces me termino poniendo el mismo pantalón negro de siempre y solo le cambio la parte de arriba.” — <i>With all the clothes I have, I don't know what to wear, so I end up in the same black trousers and only change the top.</i> It took her 20 minutes.
What she already owns	“Un vestido hermoso en Mango hace 1 mes, no lo he usado... y además luego se me olvida.” — <i>A beautiful dress from Mango a month ago; I haven't worn it... and then I forget about it.</i>
Photographing clothes	“Puede ser que después de la prenda 40 ya sea muy tedioso.” — <i>After the 40th item it might get too tedious.</i>
Privacy	“Evitaría que saliera mi cara. Con cara lo pensaría 2 veces, por miedo a que las fotos salgan.” — <i>I'd keep my face out. With my face I'd think twice, for fear the photos get out.</i>
Paying	She has never used a closet app. She pays for Disney+ every month (339 MXN).
Using /core	She did it alone, without help.

What it changed

- **Try-on has to accept a photo without a face.** She would upload a full-body photo, but not with her face.
- **The photographing threshold is not 10 items.** Ten is still the minimum to give value; 40 is where it gets tedious for her.
- **The one-time price is unproven.** She already pays a monthly subscription, and this conversation did not test whether she would pay for ClosetAI. That is the next conversation — with someone aged 18 to 35.

The quotes on the page were checked one by one against her answers before being published.

Validación con una persona real

Los datos dicen qué hace el mercado. Una conversación dice qué hace una persona.

Mujer, 57 años, Ciudad de México, organizadora de eventos. Familiar de la fundadora y fuera del segmento núcleo (18 a 35). Entrevista en persona · 2026-09-21

“Entre tanta ropa que tengo no sé qué ponerme, entonces me termino poniendo el mismo pantalón negro de siempre y solo le cambio la parte de arriba.”

“Un vestido hermoso en Mango hace 1 mes, no lo he usado por el frío o porque no he salido casi de mi casa, y además luego se me olvida.”

“Puede ser que después de la prenda 40 ya sea muy tedioso.”

“Evitaría que saliera mi cara. Con cara lo pensaría 2 veces, por miedo a que las fotos salgan.”

CONFIRMÓ

Que puede ahorrar tiempo y ayudarle a saber qué tiene en su clóset: hoy tarda 20 minutos en decidir y termina con lo mismo de siempre. Usó /core sola y sin dificultad.

CONTRADIJO

No le gusta la idea de subir su cara a internet: subiría una foto de cuerpo completo solo sin la cara. Y fotografiar la ropa cansa: después de unas 40 prendas se vuelve tedioso.

How it appears on closetai.lat/research. The summary panel now reads “1 real conversation”; the end-to-end test for this section checks that it shows the conversation and no longer says “Pending”.

12 · Iteration log — what changed after testing

WHAT TESTING FOUND	WHAT CHANGED
The search engine summarized two figures wrong	Every data point verified at its source; what could not be opened was discarded
The migration had 6 backticks that would break last week's method of applying it	Sent as escaped JSON; checked identical (13,547 characters) before running; verified against the database after
A 503 from Gemini told the person to rephrase their question	Provider failures classified (429 vs 503) in the retry and in <code>/core</code> and <code>/research</code> — last week's defect with another error code
Retrying a saturated model made people wait: measured 503, 503, and a 200 after 86 s ; the lighter model answered in 0.5 s	25 s per attempt, then the lighter model — only on provider failures. The model that answered is recorded
"22 pass, 0 fail" — but the 4 skipped were exactly the ones testing the fix	The quota gate checks both models; the 4 now run and pass. A green with skips is not a green
The report over-generalized "none of the ones reviewed" into "there are none" — while its own "to validate" line admitted the doubt	Prompt v2: keep the limit of the data. Same question, fixed
The 11th report from one IP was blocked	The per-IP limit works in production; tests now skip with that reason instead of failing
Every new account was born with the city mis-encoded ("Ciudad de Mv/©xico") since Week 0 — found in the screenshot of an unrelated flaky test	Migration 007 fixes the database default and repairs the rows; verified byte by byte; a new test creates a fresh account and checks it. In Week 1 I had fixed the symptom by hand, not the cause


One correction of my own along the way: my first script to verify prompt v2 read the question instead of the answer and reported "does not generalize ✓". A check that passes without looking at what it claims to check. It was repeated reading the answer from the database.

13 · Demo video

2:30, recorded against the live site, hosted on the product's own domain: closetai.lat/demo/semana-2, narrated by me.

Semana 2 · ¿El problema es real?

Recorrido de 2:30 por /research en producción: cada dato con su fuente, una pregunta respondida en vivo y una conversación real.



Qué muestra

- La página pública /research, sin cuenta
- El panel de resumen, calculado desde la base de datos
- Una pregunta en vivo y un informe que solo cita fuentes verificadas
- Los 5 referentes globales y la sección de México
- La tabla de competidores y sustitutos, con filtro y búsqueda sin acce...

The video page. Click it to play the video.

In order: the research page with no account; the summary panel; a question asked live and answered in seconds, citing only verified sources and saying “none of the ones reviewed”; the 5 global examples; Mexico; the table with the type filter and an accent-free search; the risk map; the real conversation; what was discarded; and the GitHub history, with the packet committed before the code.

14 · Human Decision Note

This week I had to prove the problem is real, so I started by checking my own research. I opened every source instead of trusting a search summary, and nine claims in my master document did not survive. The one that hurt most was my thesis: I believed there was no closet app in Spanish, and there are several. The gap is not the language; it is Mexico — sizes, stores and prices in pesos. I also found an app already called "Closet AI".

The decision I defend is that the AI can only reason over data a person verified. I rejected letting it search the internet: in this same research, the search engine got two figures wrong. And when a report said "there are no local competitors" while its source said "none of the ones reviewed", I rewrote the prompt.

My real conversation was with my mother, 57 — outside my 18–35 segment, and the page says so. She still changed two things: she would photograph up to 40 garments, not the 10 I assumed, and she would upload a full-body photo only without her face, so try-on has to accept that. She pays monthly for Disney+, so my belief that people only accept one-time payments is unproven. That is my next conversation.

Testing also found a bug from Week 0 that I had once fixed by hand without looking for its cause. This time I fixed the cause.

240 words · approved by Tamara Muñoz Delgadillo, September 21, 2026

Build gates

GATE	REQUIRED	EVIDENCE
1 — Plan	Build Discipline Packet complete before coding	§2 · e3f65d3 , 14 minutes before the first code commit
2 — Build	Code, commits and deploys completed	§6 · 15 commits · §7 · 14 production deployments
3 — Test	Tests documented, bug fixed, redeploy completed	§10–12 · 5 defects found by testing, each fixed and redeployed
4 — Explain	Demo and human decision note submitted	§13 · 2:30 video · §14 · 240 words